

Snowflake Streamlit Deployment Guide

Prerequisites

1. Snowflake account with appropriate permissions
2. Snowflake CLI (`snow`) installed and configured
3. Model files and data files ready for upload

Step 1: Install Snowflake CLI

```
bash

# Install Snowflake CLI
pip install snowflake-cli-labs

# Configure connection
snow connection add

# Follow prompts to enter:
# - Connection name: my_snowflake_connection
# - Account identifier: <YOUR_ACCOUNT>.snowflakecomputing.com
# - Username: <YOUR_USERNAME>
# - Password: <YOUR_PASSWORD>
# - Role: <YOUR_ROLE>
# - Warehouse: SHAP_MAP_WH
# - Database: SHAP_MAP_DB
# - Schema: GEOSPATIAL_ML

# Test connection
snow connection test --connection my_snowflake_connection
```

Step 2: Create Snowflake Resources

```
sql
```

```
-- Connect to Snowflake using SnowSQL or Snowsight
```

```
snowsql -a <ACCOUNT> -u <USERNAME>
```

```
-- Create resources
```

```
USE ROLE ACCOUNTADMIN; -- Or your admin role
```

```
CREATE DATABASE IF NOT EXISTS SHAP_MAP_DB;
```

```
CREATE SCHEMA IF NOT EXISTS SHAP_MAP_DB.GEOSPATIAL_ML;
```

```
CREATE WAREHOUSE IF NOT EXISTS SHAP_MAP_WH
```

```
WITH WAREHOUSE_SIZE = 'MEDIUM'
```

```
AUTO_SUSPEND = 300
```

```
AUTO_RESUME = TRUE
```

```
INITIALLY_SUSPENDED = TRUE;
```

```
CREATE STAGE IF NOT EXISTS SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS
```

```
DIRECTORY = (ENABLE = TRUE);
```

```
-- Grant permissions (replace <YOUR_ROLE> with your role name)
```

```
GRANT USAGE ON WAREHOUSE SHAP_MAP_WH TO ROLE <YOUR_ROLE>;
```

```
GRANT USAGE ON DATABASE SHAP_MAP_DB TO ROLE <YOUR_ROLE>;
```

```
GRANT USAGE ON SCHEMA SHAP_MAP_DB.GEOSPATIAL_ML TO ROLE <YOUR_ROLE>;
```

```
GRANT ALL ON STAGE SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS TO ROLE <YOUR_ROLE>;
```

```
-- Verify stage creation
```

```
LIST @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS;
```

Step 3: Upload Files to Stage

```
bash
```

```
# Navigate to your project directory with model and data files
cd /path/to/your/project

# Upload equipment model files
snow stage put \
    xgboost_model_v3.json \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
    --connection my_snowflake_connection \
    --overwrite

snow stage put \
    ocp_wiredown_model_data.gpkg \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/ \
    --connection my_snowflake_connection \
    --overwrite

snow stage put \
    feature_names.json \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
    --connection my_snowflake_connection \
    --overwrite

# Upload transformer model files (if you have them)
snow stage put \
    transformer_model.bst \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
    --connection my_snowflake_connection \
    --overwrite

snow stage put \
    transformer_model_data.gpkg \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/ \
    --connection my_snowflake_connection \
    --overwrite

snow stage put \
    structure_feature_names.json \
    @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
    --connection my_snowflake_connection \
    --overwrite

# Verify uploads
```

```
snow stage list @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/\  
--connection my_snowflake_connection
```

Step 4: Prepare Streamlit Application

Create the following directory structure:

```
my_shap_map_app/  
|   └── streamlit_app.py      # The complete code from the artifact  
|   └── environment.yml       # Conda dependencies  
└── .streamlit/  
    └── secrets.toml          # Snowflake credentials (for local testing)
```

environment.yml:

```
yaml  
  
name: shap_map_env  
channels:  
  - snowflake  
  - conda-forge  
dependencies:  
  - python=3.9  
  - snowflake-snowpark-python  
  - streamlit>=1.28.0  
  - pandas>=1.5.0  
  - numpy>=1.23.0  
  - geopandas>=0.12.0  
  - shapely>=2.0.0  
  - fiona>=1.9.0  
  - xgboost>=1.7.0  
  - shap>=0.42.0  
  - matplotlib>=3.6.0  
  - pyproj>=3.4.0
```

.streamlit/secrets.toml (for local testing only):

```
toml
```

```
[snowflake]
account = "<YOUR_ACCOUNT>"
user = "<YOUR_USERNAME>"
password = "<YOUR_PASSWORD>"
role = "<YOUR_ROLE>"
warehouse = "SHAP_MAP_WH"
database = "SHAP_MAP_DB"
schema = "GEOSPATIAL_ML"
```

Step 5: Test Locally (Optional)

```
bash

# Install dependencies
conda env create -f environment.yml
conda activate shap_map_env

# Run Streamlit locally
streamlit run streamlit_app.py

# Access at http://localhost:8501
```

Step 6: Deploy to Snowflake

```
bash
```

```

# Navigate to app directory
cd my_shap_map_app

# Deploy to Snowflake
snow streamlit deploy \
--connection my_snowflake_connection \
--warehouse SHAP_MAP_WH \
--database SHAP_MAP_DB \
--schema GEOSPATIAL_ML \
--replace

# Note: The app will be named based on your directory name
# You can specify a custom name with --name flag:
snow streamlit deploy \
--name SHAP_MAP_APP \
--connection my_snowflake_connection \
--warehouse SHAP_MAP_WH \
--database SHAP_MAP_DB \
--schema GEOSPATIAL_ML \
--replace

```

Step 7: Access Your Application

Option 1: Via Snowsight

1. Log in to Snowsight (<https://app.snowflake.com>)
2. Navigate to **Data > Databases**
3. Select **SHAP_MAP_DB** > **GEOSPATIAL_ML**
4. Click on **Streamlit** in the left sidebar
5. Click on your app name (e.g., **SHAP_MAP_APP**)
6. The application will launch in a new tab

Option 2: Via Direct URL

The URL format is:

`https://app.snowflake.com/<ORGNAME>/<ACCOUNT>/#/streamlit-apps/SHAP_MAP_DB.GEOSPATIAL_ML.SHAP_MAP_APP`

Step 8: Grant Access to Users

```
sql  
  
-- Grant access to other users/roles  
USE ROLE ACCOUNTADMIN;  
  
-- Grant usage on database and schema  
GRANT USAGE ON DATABASE SHAP_MAP_DB TO ROLE <TARGET_ROLE>;  
GRANT USAGE ON SCHEMA SHAP_MAP_DB.GEOSPATIAL_ML TO ROLE <TARGET_ROLE>;  
  
-- Grant access to the Streamlit app  
GRANT USAGE ON STREAMLIT SHAP_MAP_DB.GEOSPATIAL_ML.SHAP_MAP_APP TO ROLE <TARGET_ROLE>;  
  
-- Grant warehouse usage for compute  
GRANT USAGE ON WAREHOUSE SHAP_MAP_WH TO ROLE <TARGET_ROLE>;  
  
-- Grant read access to the stage  
GRANT READ ON STAGE SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS TO ROLE <TARGET_ROLE>;
```

Step 9: Update and Redeploy

When you need to update the application:

```
bash  
  
# Make changes to streamlit_app.py  
  
# Redeploy (use --replace to overwrite existing app)  
snow streamlit deploy \  
  --name SHAP_MAP_APP \  
  --connection my_snowflake_connection \  
  --warehouse SHAP_MAP_WH \  
  --database SHAP_MAP_DB \  
  --schema GEOSPATIAL_ML \  
  --replace
```

Troubleshooting

Issue: "Module not found" errors

Solution: Ensure all dependencies are in `environment.yml`:

```
yaml
```

dependencies:

- python=3.9
- snowflake-snowpark-python
- streamlit
- pandas
- numpy
- geopandas
- shapely
- fiona
- xgboost
- shap
- matplotlib
- pyproj

Issue: "File not found" when loading from stage

Solution: Verify stage paths and file existence:

```
sql
```

-- List all files in stage

```
LIST @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS;
```

-- Check specific directories

```
LIST @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/;
```

```
LIST @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/;
```

Issue: Application is slow

Solutions:

1. Increase warehouse size:

```
sql
```

```
ALTER WAREHOUSE SHAP_MAP_WH SET WAREHOUSE_SIZE = 'LARGE';
```

2. Optimize data loading with caching:

- The `@st.cache_resource` decorator is already used
- Ensure it's not being cleared unnecessarily

3. Reduce viewport limit in sidebar settings

Issue: SHAP plots fail to generate

Solution: Check that:

1. Model file is compatible with xgboost version
2. Feature names match exactly between model and data
3. There's enough memory (increase warehouse size)

Issue: Geometry rendering problems

Solution:

1. Verify GeoPackage CRS is EPSG:4326 or can be reprojected
2. Check for null/invalid geometries in data:

```
python  
  
#Add to streamlit_app.py for debugging  
st.write(f"Null geometries: {gdf.geometry.isna().sum()}"  
st.write(f"Invalid geometries: {(~gdf.geometry.is_valid).sum()}")
```

Performance Optimization

1. Warehouse Sizing

Start with MEDIUM, scale up as needed:

```
sql  
  
-- For development/testing  
ALTER WAREHOUSE SHAP_MAP_WH SET WAREHOUSE_SIZE = 'SMALL';  
  
-- For production with multiple users  
ALTER WAREHOUSE SHAP_MAP_WH SET WAREHOUSE_SIZE = 'LARGE';  
  
-- For heavy SHAP computations  
ALTER WAREHOUSE SHAP_MAP_WH SET WAREHOUSE_SIZE = 'XLARGE';
```

2. Data Optimization

Optimize GeoPackage files before uploading:

```

python

# Pre-processing script (run before upload)
import geopandas as gpd

# Load data
gdf = gpd.read_file("your_data.gpkg")

# Remove unnecessary columns
keep_cols = ['ID', 'geometry', 'feature1', 'feature2', ...] # Only what you need
gdf = gdf[keep_cols]

# Simplify geometries if too complex
gdf['geometry'] = gdf.geometry.simplify(tolerance=0.0001, preserve_topology=True)

# Ensure EPSG:4326
gdf = gdf.to_crs(epsg=4326)

# Save optimized version
gdf.to_file("optimized_data.gpkg", driver="GPKG")

```

3. Caching Strategy

The application uses `@st.cache_resource` for model loading. Monitor cache behavior:

```

python

# Add to streamlit_app.py for monitoring
st.sidebar.markdown("---")
st.sidebar.markdown("### 🔎 Cache Info")
if st.sidebar.button("Clear Cache"):
    st.cache_resource.clear()
    st.rerun()

```

4. Viewport Loading

Adjust `LIMIT` based on zoom level (advanced):

```

python

# Add to streamlit_app.py
zoom_level = st.session_state.get('map_zoom', 10)
dynamic_limit = min(5000, int(1000 * (15 - zoom_level))) # More features at higher zoom

```

Monitoring and Logging

View Streamlit Logs

```
sql  
  
-- Query Streamlit app logs  
SELECT *  
FROM TABLE(INFORMATION_SCHEMA.STREAMLIT_LOGS(  
    STREAMLIT_NAME => 'SHAP_MAP_DB.GEOSPATIAL_ML.SHAP_MAP_APP',  
    START_TIME => DATEADD('hour', -1, CURRENT_TIMESTAMP())  
))  
ORDER BY TIMESTAMP DESC;
```

Monitor Warehouse Usage

```
sql  
  
-- Check warehouse credit usage  
SELECT  
    WAREHOUSE_NAME,  
    SUM(CREDITS_USED) as TOTAL_CREDITS,  
    COUNT(*) as QUERY_COUNT  
FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY  
WHERE WAREHOUSE_NAME = 'SHAP_MAP_WH'  
    AND START_TIME >= DATEADD('day', -7, CURRENT_TIMESTAMP())  
GROUP BY WAREHOUSE_NAME;  
  
-- Check query performance  
SELECT  
    QUERY_ID,  
    QUERY_TEXT,  
    EXECUTION_TIME,  
    WAREHOUSE_SIZE  
FROM SNOWFLAKE.ACCOUNT_USAGE.QUERY_HISTORY  
WHERE WAREHOUSE_NAME = 'SHAP_MAP_WH'  
    AND START_TIME >= DATEADD('hour', -1, CURRENT_TIMESTAMP())  
ORDER BY EXECUTION_TIME DESC  
LIMIT 10;
```

Advanced Configuration

Multi-Environment Setup

Create separate environments for dev/staging/prod:

```
sql

-- Development
CREATE DATABASE SHAP_MAP_DEV;
CREATE SCHEMA SHAP_MAP_DEV.GEOSPATIAL_ML;
CREATE STAGE SHAP_MAP_DEV.GEOSPATIAL_ML.ML_ASSETS;
CREATE WAREHOUSE SHAP_MAP_DEV_WH WITH WAREHOUSE_SIZE = 'XSMALL';

-- Production
CREATE DATABASE SHAP_MAP_PROD;
CREATE SCHEMA SHAP_MAP_PROD.GEOSPATIAL_ML;
CREATE STAGE SHAP_MAP_PROD.GEOSPATIAL_ML.ML_ASSETS;
CREATE WAREHOUSE SHAP_MAP_PROD_WH WITH WAREHOUSE_SIZE = 'LARGE';
```

Deploy to different environments:

```
bash

# Deploy to dev
snow streamlit deploy \
--name SHAP_MAP_APP_DEV \
--database SHAP_MAP_DEV \
--schema GEOSPATIAL_ML \
--warehouse SHAP_MAP_DEV_WH \
--connection my_snowflake_connection

# Deploy to prod
snow streamlit deploy \
--name SHAP_MAP_APP \
--database SHAP_MAP_PROD \
--schema GEOSPATIAL_ML \
--warehouse SHAP_MAP_PROD_WH \
--connection my_snowflake_connection
```

Using Snowflake Tables Instead of GeoPackage

For very large datasets, consider loading data into Snowflake tables:

```

sql

-- Create table for equipment data
CREATE TABLE SHAP_MAP_DB.GEOSPATIAL_ML.EQUIPMENT_DATA (
    ID VARCHAR,
    HAS_WIREDOWN INTEGER,
    POF FLOAT,
    GEOMETRY GEOGRAPHY,
    -- Add all your feature columns
    FEATURE1 FLOAT,
    FEATURE2 FLOAT,
    -- ... more features
);
-- Load data from stage (after converting to CSV/Parquet)
COPY INTO SHAP_MAP_DB.GEOSPATIAL_ML.EQUIPMENT_DATA
FROM @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/equipment_data.parquet
FILE_FORMAT = (TYPE = PARQUET);

```

Then modify the loading code in `streamlit_app.py`:

```

python

# Replace load_geopackage_from_stage with:
def load_data_from_table(table_name: str) -> gpd.GeoDataFrame:
    """Load data from Snowflake table"""
    df = session.table(table_name).to_pandas()

    # Convert GEOGRAPHY to shapely geometries
    from shapely import wkb
    df['geometry'] = df['GEOMETRY'].apply(lambda x: wkb.loads(bytes.fromhex(x)))

    gdf = gpd.GeoDataFrame(df, geometry='geometry', crs='EPSG:4326')
    return gdf

```

Security Best Practices

1. Role-Based Access Control

```

sql

```

```
-- Create custom roles
CREATE ROLE SHAP_MAP_VIEWER;
CREATE ROLE SHAP_MAP_ADMIN;

-- Grant viewer permissions (read-only)
GRANT USAGE ON DATABASE SHAP_MAP_DB TO ROLE SHAP_MAP_VIEWER;
GRANT USAGE ON SCHEMA SHAP_MAP_DB.GEOSPATIAL_ML TO ROLE SHAP_MAP_VIEWER;
GRANT USAGE ON STREAMLIT SHAP_MAP_DB.GEOSPATIAL_ML.SHAP_MAP_APP TO ROLE SHAP_MAP_VIEWER;
GRANT USAGE ON WAREHOUSE SHAP_MAP_WH TO ROLE SHAP_MAP_VIEWER;
GRANT READ ON STAGE SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS TO ROLE SHAP_MAP_VIEWER;

-- Grant admin permissions (can update models/data)
GRANT ALL ON DATABASE SHAP_MAP_DB TO ROLE SHAP_MAP_ADMIN;
GRANT ALL ON SCHEMA SHAP_MAP_DB.GEOSPATIAL_ML TO ROLE SHAP_MAP_ADMIN;
GRANT ALL ON STAGE SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS TO ROLE SHAP_MAP_ADMIN;

-- Assign roles to users
GRANT ROLE SHAP_MAP_VIEWER TO USER data_analyst_user;
GRANT ROLE SHAP_MAP_ADMIN TO USER ml_engineer_user;
```

2. Secrets Management

Never hardcode credentials in `streamlit_app.py`. Always use Snowflake secrets:

```
python

# Good - using Snowflake connection
session = st.connection("snowflake").session()

# Bad - hardcoding credentials
# connection_parameters = {
#     "account": "myaccount",
#     "password": "hardcoded_password" # NEVER DO THIS
# }
```

3. Network Policies

sql

```
-- Restrict access to specific IP ranges
CREATE NETWORK POLICY SHAP_MAP_POLICY
ALLOWED_IP_LIST = ('192.168.1.0/24', '10.0.0.0/8');

ALTER USER data_analyst_user SET NETWORK_POLICY = SHAP_MAP_POLICY;
```

Maintenance Tasks

Regular Updates

Monthly:

1. Check warehouse credit usage
2. Review query performance logs
3. Update model files if retrained
4. Check for Streamlit/dependency updates

Quarterly:

1. Review user access permissions
2. Optimize data storage (remove old versions)
3. Update documentation
4. Evaluate warehouse sizing

Model Updates

To update models without redeploying the entire app:

```
bash

# Upload new model version
snow stage put new_model_v4.json \
@SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
--overwrite

# Update MODEL_CONFIGS in streamlit_app.py to point to new file
# Then redeploy
snow streamlit deploy --name SHAP_MAP_APP --replace
```

Data Refresh

```
bash

# Upload updated data
snow stage put updated_data.gpkg \
@SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/ \
--overwrite

# Clear cache in app to reload data
# (Add cache clear button in sidebar as shown above)
```

Cost Management

1. Auto-Suspend Warehouses

```
sql

ALTER WAREHOUSE SHAP_MAP_WH SET
  AUTO_SUSPEND = 60 -- Suspend after 1 minute of inactivity
  AUTO_RESUME = TRUE;
```

2. Set Resource Monitors

```
sql

CREATE RESOURCE MONITOR SHAP_MAP_MONITOR
  WITH CREDIT_QUOTA = 100 -- 100 credits per month
    FREQUENCY = MONTHLY
    START_TIMESTAMP = IMMEDIATELY
    TRIGGERS
      ON 75 PERCENT DO NOTIFY
      ON 100 PERCENT DO SUSPEND
      ON 110 PERCENT DO SUSPEND_IMMEDIATE;

ALTER WAREHOUSE SHAP_MAP_WH SET RESOURCE_MONITOR = SHAP_MAP_MONITOR;
```

3. Monitor Costs

```
sql
```

```
-- Daily credit usage
SELECT
    DATE(START_TIME) as DATE,
    SUM(CREDITS_USED) as DAILY_CREDITS
FROM SNOWFLAKE.ACCOUNT_USAGE.WAREHOUSE_METERING_HISTORY
WHERE WAREHOUSE_NAME = 'SHAP_MAP_WH'
    AND START_TIME >= DATEADD('day', -30, CURRENT_TIMESTAMP())
GROUP BY DATE(START_TIME)
ORDER BY DATE DESC;
```

Backup and Recovery

Backup Models and Data

```
bash

# Download files from stage to local backup
snow stage get @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/ \
/local/backup/models/ \
--connection my_snowflake_connection

snow stage get @SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/data/ \
/local/backup/data/ \
--connection my_snowflake_connection
```

Version Control

Keep versions of your models:

```
bash

# Upload with version suffix
snow stage put xgboost_model_v3.json \
@SHAP_MAP_DB.GEOSPATIAL_ML.ML_ASSETS/models/archive/xgboost_model_v3_2025_01.json
```

Additional Resources

- **Snowflake Streamlit Documentation:** <https://docs.snowflake.com/en/developer-guide/streamlit/about-streamlit>
- **Snowflake CLI Documentation:** <https://docs.snowflake.com/en/developer-guide/snowflake-cli/index>
- **Streamlit Documentation:** <https://docs.streamlit.io/>

- **XGBoost Documentation:** <https://xgboost.readthedocs.io/>
- **SHAP Documentation:** <https://shap.readthedocs.io/>
- **Geopandas Documentation:** <https://geopandas.org/>

Support and Community

- **Snowflake Community:** <https://community.snowflake.com/>
- **Streamlit Community:** <https://discuss.streamlit.io/>
- **GitHub Issues:** Report issues specific to your deployment

Checklist for Go-Live

- All resources created (database, schema, warehouse, stage)
- Files uploaded to stage and verified
- Permissions granted to all users/roles
- Application deployed and tested
- Resource monitors configured
- Auto-suspend enabled on warehouse
- Documentation shared with users
- Backup procedures established
- Monitoring queries saved
- Support contacts identified

Note: Replace all placeholder values (`<YOUR_ACCOUNT>`, `<YOUR_ROLE>`, etc.) with your actual Snowflake account details before executing commands.